Automated Identification and Extraction of Op-Ed Predictions Using Machine Learning[*]

Lawrence Chan
University of Pennsylvania
Philadelphia, PA 19104
`chanlaw@seas.upenn.edu`

Abstract

By scoring the predictive accuracy of political pundits, we can hold them accountable for bad predictions and reward better pundits. However, in order to score the predictive accuracy of pundits, their predictions must be first identified and extracted by hand, a laborious process that is a key limiting factor of this endeavor. In this paper we propose to use machine learning techniques to automate the identification and extraction process. We train various machine learning models and show that this approach has much promise.

1 Introduction

The problem of poor predictive accuracy of the average pundits is well known amongst the forecasting literature. For example, a study of 284 experts and over 28,000 predictions, conducted by Philip Tetlock over a period of more than 20 years, found that experts performed barely better than chance (Tetlock, 2005). However, research has also shown a large degree of variation among the accuracy of forecasts, and so there is significant interest in tracking and scoring the accuracy of political pundits (Tetlock, 2016). Many even hope to make this a standard part of standard intellectual discourse, much like polling numbers are a necessary part of any discussion about election forecasts today.

Nevertheless, this process of identifying and extracting predictions from the vast volume of editorial

Prediction	Non-prediction
Trump is a long shot for the general election.	Trump is a con man and a crook.
Returning to the gold standard will cause a recession.	We should increase funding for our children's education.
Sanders will be a stronger general election candidate than Clinton.	Clinton's dominance was based on her lead among non-whites.

Table 1: Examples of predictions and non-predictions

pieces published daily is a daunting task, even if we were to limit our attention to only a few of the most notable pundits. In this paper we propose using machine learning to partially resolve this problem. By trawling the writings of these pundits and returning excerpts containing their predictions, we can vastly reduce the amount of manpower needed for such a project.

1.1 What we mean by prediction

In this context, a prediction is a statement that constrains our beliefs about the future. Predictions need not be certain - a statement that a certain future event is very likely or unlikely is also a prediction about the future, especially if the speaker assigns a probability estimate to the prediction. In order to be scored, predictions should be about events that we can verify objectively, at least in principle, a condition known as the "Clairvoyance Test" (Tetlock and Gardner, 2015). Examples of sentences that we consider predictions and non-predictions are included in **Table 1**.

Articles	Sentences	Unique Words
207	3882	7621

Table 2: Summary of the dataset used. Articles were collected from Paul Krugman's New York Times Column.

2 Approach

For this project, we decided on the following three-step approach for training a classifier to recognize predictions:

1. **Webscraping and Data Cleaning** First, we used webscraping tools to acquire 207 articles from the New York Times Column of Paul Krugman, a nobel-prize winning economist and one of the most famous liberal pundits in America. We then cleaned the dataset and segmented it into 3882 sentences.

2. **Crowdsourcing** We then used Crowdflower, a crowdsourcing website, to label the 3882 sentences as containing predictions or not containing predictions.

3. **Machine Learning Classifier** Finally, we vectorized the sentences using a bag of words model, and then trained various machine learning models on this dataset, using a 80% training, 10% validation, and 10% test set split.

More explanation for the rationale behind this three step approach is detailed in the next subsection, and more specifics about each of the three steps are available in each of their corresponding sections. A summary of the dataset is presented in table 2. The code and data used in this project is available online. [1]

2.1 Rationale

Our first attempt at classifying predictions was to use Hearst-like syntatical patterns to identify predictions (Hearst, 1992). However, despite creating combinations of rules, some with up to a dozen rules, we found that we could not beat the rule that always predicted that the given sentence contains no predictions on a hand-gathered test set of 150 sentences. This is in line with research showing the limitations of Hearst patterns in comparison to machine learning based models (Snow et al., 2004). As such, we

---

[1] `github.com/chanlaw/prediction-extraction`

abandoned this approach in favor of one using machine learning classifiers.

Since no corpus of prediction-containing sentences currently exists, we decided to create our own dataset using webscraping and crowdsourcing. After considering various political pundits, we settled on Paul Krugman as a famous pundit who made significant more predictions than other pundits, averaging more than one per article. (For example, the Pulitzer-winning journalist Thomas Friedman makes almost no concrete predictions at all in his articles.) Because of the size of the dataset required to train a classifier, we decided to label the dataset using crowdsourcing.

We explored segmenting the dataset using other methods - both in terms of longer excerpts (at paragraph level) and in terms of shorter excerpts (at the phrase level). Ultimately, we chose to segment the articles at the sentence-level, because this seemed like a good compromise between having sufficiently a large dataset and having enough information to identify predictions.

3 Webscraping and Data Cleaning

We used Python's `urllib2` package and the `Beautiful Soup` Python Library to extract the content from the first 55 pages of Paul Krugman's column, The Conscience of a Liberal. We took advantage of the fact that the formatting of all New York Times columns is standardized; by extracting html items with the class "story-body-text", we were able to get the content of the articles without having to use an API to clean the scrapped html. A total of 207 articles were extracted (Krugman's weekly music sharing posts were ignored). We then used NLTK, a natural language processing package, to segment the articles into sentences. We obtained a total of 3882 sentences through this process.

4 Crowdsourcing

After we obtained the dataset of segmented sentences, we used Crowdflower to label the sentences. Crowdflower was chosen at the suggestion of Professor Callison-Burch for its ease of use and its support for data validation by using test questions to filter out untrustworthy contributors. We paid contributors $0.11 for every set of 8 sentences (7 to be

Prediction Detection

Instructions ▲

Overview

Decide of the sentence contains a prediction - that is a testable statement about how the future will turn out (at the time of writing). Note that predictions need not concrete, instead, they should highlight the speaker's belief that certain outcomes are more or less likely. This means that many predictions talk about odds or likelihoods. We are looking for three types of predictions: outcome predictions, scenarios, and counterfactual comparisons.

Definitions and Examples

We are looking for the following three kinds of predictions:

- **Outcome predictions:** concrete statements about what will happen, often with an implied or explicit timeline.
  - **Examples:** "There will be no new terrorist attack against the USA in 2016." "Trump is a long shot for the general election." "The risks tilt heavily toward inflation below official targets."
- **Scenarios:** conditional predictions, claiming that the conditions of the scenario are met, then certain outcomes will follow.
  - **Examples:** "If we institute NGDP targeting, we will see a more stable economy." "Returning to the gold standard will cause a recession." "A Trump presidency will probably make America great again."
- **Counterfactual comparisons:** judgments about the relative merit of two possible worlds.
  - **Examples:** "President Trump will be better for the US economy than a President Cruz." "Bernie Sanders will be a stronger general election candidate than Hillary Clinton."

Examples of Sentences not Containing Predictions

- **Statements about people's characters.** "Trump is a con man".
- **Explanations for past events.** "Mrs Clinton's dominance on Super Tuesday, despite four victories for her rival, Bernie Sanders, in his own state of Vermont, in Oklahoma, Colorado and Minnesota, was based on her crushing lead among non-whites."
- **Factual statements about the present.** "India's small businesses have limited access to loans; most of its consumers do not have credit cards, or for that matter credit."
- **Policy suggestions without projections.** "We should increase funding for our children's educations."

Thank You!

Your careful attention on this task is greatly appreciated!

Figure 1: Final task description given to contributors.

Attempt	Overall	Instructions
Test Run 1	3.4	3.6
Test Run 2	3.1	3.6
Test Run 3	4.2	4.3
Final Run	4.2	4.3

Table 3: Contributor ratings of labelling task (out of 5)

labelled and 1 test). In order to guarantee accuracy, we had each sentence labelled by 3 different contributors, and took the majority opinion for each sentence. The total cost for all 3882 sentences, including the 20% fee charged by Crowdflower and payments to individuals who failed the test questions, was $250.95, or $0.064 per sentence and $1.21 per article.

We used an iterative process to design the Crowdsourcing instructions for the task. We wrote three different sets of instructions and ran each of them through Crowdflower with a test set of 100 sentences. The contributor ratings for the three test runs are displayed in **Table 3**. We then used the best instruction set, which we display in **Figure 1**, to label the entire dataset.

In total the crowd identified 454 predictions in the 3882 sentences.

5 Machine Learning Classifier

After removing punctuation and capitalization, we segmented each sentence into words. We then vectorized the sentences using a bag of words (unigram) model of the most frequent 1000 words, from a dataset containing 7621 unique words. We also used two alternate vectorizations, one using the 10,000 most frequent unigrams and bigrams, and another using the 50,000 most frequent unigrams, bigrams, and trigrams.

With this dataset, we turned trained four kinds of classifiers: regularized logistic regression, support vector machines (SVMs) with both RBF and linear kernels, random forests with 10 and 100 trees, and naive Bayes.

For each classifier, we performed 20-fold cross validation. In each fold, we split the dataset into 80% training, 10% validation, and 10% test set, fit the models on the training set and used the validation set to tune the regularization parameter. For our logistic regression and support vector classifiers, the regularization parameter was C, the inverse regularization strength, chosen over a grid between 0.01 and 100; for random forest, it was the max depth of the tree, which varied between 1 and 30; and for naive Bayes, it was the smoothing parameter for

Classifier	Accuracy	Recall	Precision
Logistic (50k Most Frequent Uni/Bi/Trigrams)	0.911	0.358	0.755
Logistic (10k Most Frequent Uni/Bigrams)	0.911	0.347	0.751
SVM (RBF Kernel)	0.911	0.382	0.708
Logistic	0.908	0.354	0.711
SVM (Linear Kernel)	0.907	0.377	0.681
Random Forest (100 Trees)	0.899	0.124	0.881
Naive Bayes	0.892	0.319	0.616
Random Forest (10 Trees)	0.890	0.153	0.739
Naive Bayes (50k Most Frequent Uni/Bi/Trigrams)	0.890	0.163	0.704
Naive Bayes (10k Most Frequent Uni/Bigrams)	0.886	0.179	0.674
Majority Class Classifier	0.883	0.00	N/A

Table 4: Averaged test-set performance of classifiers, in declining order of accuracy. Except where otherwise noted, models were trained with a feature set consisting of the 1000 most common unigrams.

unseen points, chosen over a grid between 0.1 and 100. We then recorded the performance of each classifier on the test set. Due to limitations on computing power, we were not able to train a support vector machine and random forest classifier on the unigram/bigram mix or the unigram/bigram/trigram mix.

6 Results

Our results are displayed in **Table 4**. Overall, logistic regression over the uni/bi/trigram mix, logistic regression over the uni/bigram mix, and SVM using the RBF kernel did the best, with an average test-set accuracy of 0.911. Logistic regression on the unigram dataset and SVM using a linear kernel were in the second tier, with an average test-set accuracy of 0.908 and 0.907, respectively. This makes sense because both logistical regression including bigrams and SVMs using the RBF kernel can draw nonlinear boundaries around words, while logistic regression using only unigrams and Linear SVMs can only draw linear boundaries with respect to the words in the feature set. This is further supported by the most informative features for the models (displayed in **Table 5**).

Note that random forest classifiers and naive Bayes performed very poorly on the dataset. The majority class classifier, which always predicts "no" on every sentence, performs almost as good as these classifiers. This is in line with the properties of the classifiers - while support vector machines and logistic regression are often used for text classifica-

Logistic (Uni/Bi/Trigram)	Logistic (Uni/Bigram)
will	will
would	wont
going to	would
if	going to
wont	if

Logistic (Unigram)	SVM (Linear Kernel)
will	will
going	going
wont	wont
raise	raise
likely	unless

Table 5: Most informative features for linear classifiers

tion, random forest classifiers and naive Bayes do much more poorly in this domain due to the assumptions they make about the data (Hastie et al., 2013).

Using our best performing logistic regression model, we acquired and labelled a new dataset of 272 sentences from Paul Krugman's column. 11 sentences were labelled as predictions. The labelled dataset was then hand-graded for accuracy, and we identified 3 false positives and 11 false negatives, meaning our model had recall 0.42 and precision 0.727. We present representative errors in **Table 6**. All three of the false positives have phrases that seem very much like real predictions made by the author, but were either questions asked by the author or commentary on predictions made by other. The false negatives were more varied; some are predictions not including words such as "will" or "would", while others were simply misclassified for no easily

False Positives	False Negatives
If Britain does leave the EU, and escapes the grip of the Eurocrats, who will it be empowering instead?	Falling oil leads to falling inflation expectations, and since interest rates can't fall, real rates go up, hurting recovery.
But what really amazes me about the latest set of stories is the promise that Ryan will finally deliver the Republican Obamacare alternative that his colleagues in Congress have somehow failed to produce after all these years.	True socialized medicine an American NHS would be feasible economically; so would single-payer, in the form of Medicare for all.

Table 6: Representative errors on new dataset

identified reason.

7 Discussion

We believe that the results above demonstrate that this approach is a viable one. By using webscraping and machine learning techniques, we can significant reduce the amount of time used in labeling extracting predictions, saving manpower for other uses. In particular, the high precision of our classifiers could greatly reduce the amount of text a researcher needs to read to identify a prediction.

However, we believe that there exists significant room for improvement. Even an accuracy rate of 0.911, the best we could achieve, is not particularly impressive, given that the majority class classifier achieves an accuracy of 0.883. First, the classifiers we learnt were all relatively low recall - the highest recall we achieved was less than 0.4. By labelling and acquiring a larger dataset and training classifiers on that dataset, we can hopefully learn better classifiers. Because the cost of crowdsourcing is relatively low, this should be an cost-effective method for getting a better classifier; another 10000 sentences will cost less than $700.

Another area for improvement is finding a better source of labels. The crowdsourced labels could probably be higher quality - the average agreement between crowd-sourced labels and hand-labelled dataset is around 95%. We can do this by either paying more and recruiting higher quality crowdworkers, or by recruiting Penn students to label the data.

References

[Hastie et al.2013] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. 2013. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction.* Springer, 10 edition.

[Hearst1992] Marti A. Hearst. 1992. Automatic acquisition of hyponyms from large text corpora. In *Proceedings of COLING 1992*, pages 23–28, Nantes, France. Association for Computational Linguistics.

[Snow et al.2004] Rion Snow, Daniel Jurafsky, and Andrew Y. Ng. 2004. Learning syntactic patterns for automatic hypernym discovery. Vancouver, Canada.

[Tetlock and Gardner2015] Philip E. Tetlock and Dan Gardner. 2015. *Superforecasting: The Art and Science of Prediction.* Crown Publishers.

[Tetlock2005] Philip E. Tetlock. 2005. *Expert Political Judgment: How Good is It? How Can we Know?* Princeton University Press, Princeton, NJ.

[Tetlock2016] Philip E. Tetlock. 2016. Revolutionizing the assessment of alpha pundits: the superforecaster challenge. Proposal, Wharton School, University of Pennsylvania.